

Class09: Candy Mini Project

Sofia Jaravata (A19160915)

Table of contents

Background	2
Data Import: Importing Candy Data	2
Q1. How many different candy types are in this dataset?	3
Q2. How many fruity candy types are in the dataset?	3
2.2 What is your favorite candy?	4
Q3. What is your favorite candy (other than Twix) in the dataset and what is its winpercent value?	4
Q4. What is the winpercent value for “Kit Kat”?	4
Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?	5
Side-note: the skimr::skim() function	5
Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?	6
Q7. What do you think a zero and one represent for the <code>candy\$chocolate</code> column?	6
3. Exploratory Analysis	6
Q8. Plot a histogram of <code>winpercent</code> values using both base R and ggplot2. . .	6
Q9. Is the distribution of <code>winpercent</code> values symmetrical?	8
Q10. Is the center of the distribution above or below 50%?	8
Q11. On average is chocolate candy higher or lower ranked than fruit candy? .	8
Q12. Is this difference statistically significant?	9
4. Overall Candy Rankings	10
Q13. What are the five least liked candy types in this set?	10
Q14. What are the top 5 all time favorite candy types out of this set?	10
Q15. Make a first barplot of candy ranking based on winpercent values.	11
Q16. This is quite ugly, use the <code>reorder()</code> function to get the bars sorted by winpercent?	13

4.0.1 Time to add some useful color	15
Q17. What is the worst ranked chocolate candy?	17
Q18. What is the best ranked fruity candy?	17
5. Taking a look at pricepercent	17
Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?	18
Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?	18
6. Exploring the correlation structure	19
Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?	19
Q23. Use your corplot result to make a prediction: which variables do you expect will have the largest contributions (i.e. loadings) to PC1 (i.e., drive the most separation between candies along the first principal component)?	20
7. Principal Component Analysis	20
Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?	25
8. Summary	25
Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?	25

Background

Today we are taking a small detour to analyze a dataset (that we have more intrinsic insight into) with the most useful analysis method we have learned thus far - Principal Component Analysis (PCA)

Data Import: Importing Candy Data

```
candy <- read.csv("candy-data.csv", row.names=1) #reads csv and organizes it
head(candy) #displays first few elements
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisp	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0

	hard	bar	pluribus	sugar	percent	price	percent	win	percent
100 Grand	0	1	0		0.732		0.860	66.97	173
3 Musketeers	0	1	0		0.604		0.511	67.60	294
One dime	0	0	0		0.011		0.116	32.26	109
One quarter	0	0	0		0.011		0.511	46.11	650
Air Heads	0	0	0		0.906		0.511	52.34	146
Almond Joy	0	1	0		0.465		0.767	50.34	755

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

85 different candy types.

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

38 fruity candy types are in the dataset.

2.2 What is your favorite candy?

```
library(dplyr)# CONSOLE: install.packages("dplyr")
```

Warning: package 'dplyr' was built under R version 4.5.2

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
candy |>
  filter(row.names(candy)=="Reese's pieces") |>
  select(winpercent)
```

```
      winpercent
Reese's pieces 73.43499
```

Q3. What is your favorite candy (other than Twix) in the dataset and what is it's winpercent value?

My favorite candy is Reese's pieces and its winpercent value is 73.43499%

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

Kit Kat's winpercent is 76.7686%.

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bar", ]$winpercent
```

```
[1] 49.6535
```

Tootsie Roll Snack Bar’s winpercent is 49.6335%.

Side-note: the `skimr::skim()` function

```
#library("skimr") #requires loading package in console
#skim(candy)

#no need to load whole package, takes one thing from skimr package efficiently
skimr::skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

From your use of the `skim()` function use the output to answer the following:

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The histogram “hist” column looks to be on a different scale to the majority of the other columns in the dataset.

Q7. What do you think a zero and one represent for the `candy$chocolate` column?


For the `candy$chocolate` column, the zero and the one is most likely a binary operator where “0” means the candy does not have chocolate and “1” means the candy has chocolate.

3. Exploratory Analysis

Q8. Plot a histogram of `winpercent` values using both base R and `ggplot2`.

```
hist(candy$winpercent)
```


Histogram of candy\$winpercent



```
library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.5.2

```
ggplot(candy) +  
  aes(winpercent)+  
  geom_histogram(bins = 20)
```



Q9. Is the distribution of winpercent values symmetrical?

No, they skew right a little bit.

Q10. Is the center of the distribution above or below 50%?

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

The center of the distribution is mostly below 50% as the median is 47.83.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Yes, chocolate candy is higher ranked than fruit candy as the mean of chocolate candy winpercent is 60.92153, while the mean of fruity candy winpercent is 44.11974.


```
#chocolate candy
choc.ind <- as.logical(candy$chocolate) #true and falses for choc candy
choc.candy <- candy[choc.ind, ] #access just choc candy
choc.win <- choc.candy$winpercent # access just winpercent values
mean(choc.win) #print out mean function of the winpercent
```

```
[1] 60.92153
```

```
#fruity candy
fruity.ind <- as.logical(candy$fruity) #true and falses for fruity candy
fruity.candy <- candy[fruity.ind, ] #access just fruity candy
fruity.win <- fruity.candy$winpercent # access just winpercent values
mean(fruity.win) #print out mean function of the winpercent
```

```
[1] 44.11974
```

Q12. Is this difference statistically significant?

```
t.test(choc.win, fruity.win)
```

Welch Two Sample t-test

```
data:  choc.win and fruity.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

This difference is statistically significant, as the p-value, 2.871e-08, is much less than 0.05 and thus statistically significant.

4. Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
#sorts the winpercent from lowest percent to highest using base R
head(candy[order(candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0

	crisped	rice	wafers	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip				0	0	0	1	0.197		0.976
Boston Baked Beans				0	0	0	1	0.313		0.511
Chiclets				0	0	0	1	0.046		0.325
Super Bubble				0	0	0	0	0.162		0.116
Jawbusters				0	1	0	1	0.093		0.511

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744

The five least liked candy types in this set is Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble, and Jawbusters.

Q14. What are the top 5 all time favorite candy types out of this set?

```
head(candy[order(candy$winpercent, decreasing = TRUE),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0

Snickers	1	0	1	1	1
	crisped	ricewafer	hard bar	pluribus	sugarpercent
Reese's Peanut Butter cup		0	0	0	0.720
Reese's Miniatures		0	0	0	0.034
Twix		1	0	1	0.546
Kit Kat		1	0	1	0.313
Snickers		0	0	1	0.546


	pricepercent	winpercent
Reese's Peanut Butter cup	0.651	84.18029
Reese's Miniatures	0.279	81.86626
Twix	0.906	81.64291
Kit Kat	0.511	76.76860
Snickers	0.651	76.67378

The top 5 all time favorite candy types are Reese's Peanut Butter Cup, Reese's Miniatures, Twix, Kit Kat, and Snickers.

Q15. Make a first barplot of candy ranking based on winpercent values.

```
library(ggplot2) #CONSOLE: install.packages("ggplot2")


ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by winpercent?

```
library(ggplot2)

ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col()
```



4.0.1 Time to add some useful color


```
my_cols=rep("black", nrow(candy))

#chocolate candy in chocolate color
my_cols[as.logical(candy$chocolate)] = "chocolate"

#candy bars in brown
my_cols[as.logical(candy$bar)] = "brown"

#fruity candy in pink
my_cols[as.logical(candy$fruity)] = "pink"

ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols) + #fills columns as different candy type colors
  ylab("")
```



Q17. What is the worst ranked chocolate candy?

The worst ranked chocolate candy is Sixlets.

Q18. What is the best ranked fruity candy?

The best ranked fruity candy is Starburst.

5. Taking a look at pricepercent


We can fix the label text overplotting with an add on package called **ggrepel** and its `geom_text_repel()`.

Make a plot of winpercent vs. pricepercent

```
library(ggrepel)# CONSOLE: install.packages("ggrepel")
```

Warning: package 'ggrepel' was built under R version 4.5.2

```
# How about a plot of win vs price
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

The candy type the highest ranked in terms of winpercent for the least money is the Reese's Miniatures because it is in the lowest right corner of the plot. It has a high winpercent and low pricepercent.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
# order price percent from high to low
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

The top 5 most expensive candy types are Nik L Nip, Nestle Smarties, Ring pop, Hershey's Krackel, Hershey's Milk Chocolate. The least popular is Nik L Nip (22.44534% winpercent).


6. Exploring the correlation structure

We can calculate the pair-wise correlation of all our columns

```
library(corrplot) # CONSOLE: install.packages("corrplot")
```

corrplot 0.95 loaded

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

The two variables “Chocolate” and “Candy” are anti-correlated because they have red values for each other.

Q23. Use your corrplot result to make a prediction: which variables do you expect will have the largest contributions (i.e. loadings) to PC1 (i.e., drive the most separation between candies along the first principal component)?

The variables “chocolate”, “fruity”, and “bar” will have the largest contributions to PC1 because they have the largest contributions.

7. Principal Component Analysis

In this case we want to be sure to set `scale=TRUE` argument for `prcomp()` because we have one variable, `winpercent`, that is on a very different scale than all others and would otherwise dominate our PCA results.

```
pca <- prcomp(candy, scale = TRUE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

First major result figure is the “score plot” of PC1 vs PC2 - how the different candy are related to each other on our new PC axis:

```
pca <- prcomp(candy, scale = TRUE)
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000


Now we can plot our main PCA score plot of PC1 vs PC2.

```
plot(pca$x[,1:2])
```



We can change the plotting character and add some color:

```
plot(pca$x[,1:2], col=my_cols, pch=16)
```




We can make a much nicer plot with the ggplot2 package but it is important to note that ggplot works best when you supply an input data.frame that includes a separate column for each of the aesthetics you would like displayed in your final plot. To accomplish this we make a new data.frame here that contains our PCA results with all the rest of our candy data. We will then use this for making plots below.

```
# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, pca$x[,1:3])
```

```
p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)
```


```
p
```



```
library(ggrepel)
p +
  geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
        subtitle="Colored by type: chocolate bar (dark brown),
        chocolate other (light brown),
        fruity (red),
        other (black)",
        caption="Data from 538")
```

Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown),
chocolate other (light brown),
fruity (red),
other (black)




Data from 538

If you want to see more candy labels you can change the max.overlaps value to allow more overlapping labels or pass the ggplot object p to plotly like so to generate an interactive plot that you can mouse over to see labels:

```
#library(plotly)
#ggplotly(p)
```

```
ggplot(pca$rotation) +
  aes(PC1, reorder(row.names(pca$rotation), PC1)) +
  geom_col()
```

Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

The variables picked up strongly by PC1 in the positive direction are “hard”, “pluribus”, and “fruity”. This makes sense because in the correlation plot earlier, those three variables showed anti-correlation, with negative values and red dots, opposing the chocolate candies.

8. Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

A combination of chocolate, bar , and not fruity characteristics appear to make a “winning” candy. The visualizations display the ranking patterns by winpercent and groups by different types of candy, the correlation matrix shows the relationships between these variables (candy elements), and the PCA separates chocolate/bar candies from fruity/pluribus/hard candies in a new axis of variation, which supports the correlation matrix structure.